



Compiler Design

Introduction to Compilers

Dr. Biagio Cosenza | TU Berlin | Wintersemester 2017-18



My Background

- PhD at the University of Salerno, Italy on Parallel Processing (2007-11)
 - HPC-Europa2 and HPC-Europa++ at HLRS Supercomputing, Stuttgart
 - ISCRA at CINECA Supercomputing, Bologna
 - DAAD Fellowship at VISUS, Universität Stuttgart
- Postdoctoral Researcher at the University of Innsbruck, Austria (2011-2015)
 - Insieme Compiler
 - Multi-disciplinary Research Platform (FWF DK-Plus Program)
- Senior Researcher at AES, TU Berlin, Germany (since April 2015)
 - Compilers and Software Optimization, High Performance Computing, Embedded Systems
 - International research projects



Lecture Organization

- Lecture
 - Tuesday, 12:00 to 14:00 in room H 2032
- Lab
 - Friday, 10:00 to 12:00 in rooms TEL 106 li and re (Telefunken-Hochhaus)
 - Teaching assistant: Nikita Popov
 - First lab is on October 27, 2017
- Final grade (Portfolio)
 - 50% final examination
 - 50% lab exercises
- Contact: (preferred) use the forum on ISIS
 - Other students may have the same problem

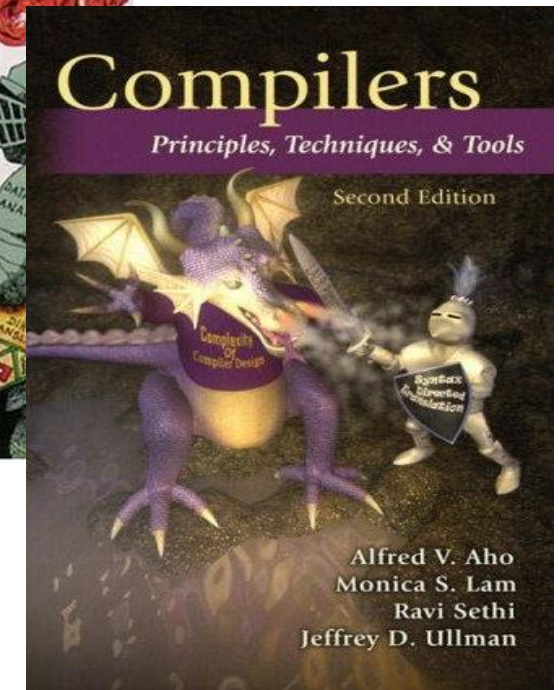
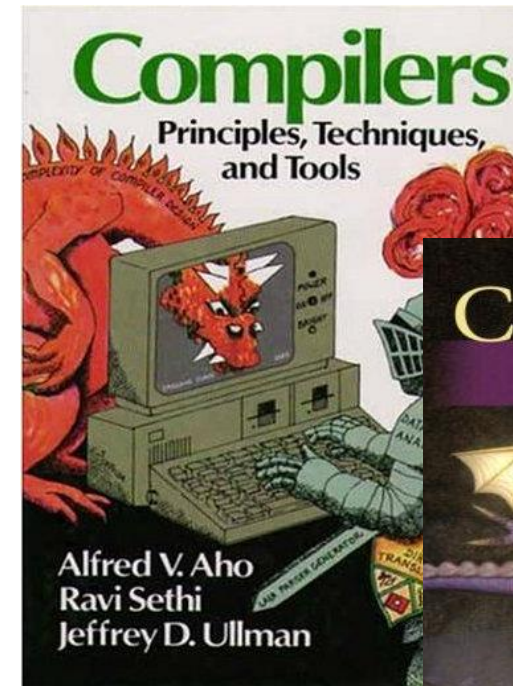
Tentative Schedule

| Week | Lecture | | Lab | |
|------|---------|---------------------------------------|--------|-----------------------------------|
| 1 | 17-Oct | Introduction | | |
| 2 | 24-Oct | Lexical Analysis | 27-Oct | P1: Lex/Flex |
| 3 | 31-Oct | Lexical Analysis | 3-Nov | P1 |
| 4 | 7-Nov | Syntax Analysis | 10-Nov | P1 |
| 5 | 14-Nov | Syntax Analysis | 17-Nov | P1 |
| 6 | 21-Nov | Syntax Analysis | 24-Nov | P2: Bison |
| 7 | 28-Nov | Semantic Analysis | 1-Dec | P2 |
| 8 | 5-Dec | Intermediate Representations | 8-Dec | P2 |
| 9 | 12-Dec | Dataflow Analysis | 15-Dec | P3: LLVM IR Analysis/Optimization |
| 10 | 19-Dec | Dataflow Analysis | 22-Dec | P3 |
| | 26-Dec | no lecture | 29-Dec | no lab |
| | 2-Jan | no lecture | 5-Jan | no lab |
| 11 | 9-Jan | SSA | 12-Jan | P3 |
| 12 | 16-Jan | Runtime Env., Code Generation | 19-Jan | P3 |
| 13 | 23-Jan | Registry Allocation | 26-Jan | P4: LLVM IR Backend |
| 14 | 30-Jan | Instruction Scheduling, Optimizations | 2-Feb | P4 |
| 15 | 6-Feb | Compilation for Embedded Systems | 9-Feb | P4 |
| 16 | 13-Feb | Exam exercise | 16-Feb | P4 |

Updated calendar available on ISIS <https://isis.tu-berlin.de/course/view.php?id=8583>

Course Text Books

- The “**Dragon Book**”, 2nd edition [ALSU]
 - Aho, Lam, Sethi, Ullman. “Compilers: Principles, Techniques and Tools”, 2nd edition
 - international version has no dragon!
 - 1st edition by Aho, Sethi, Ullman [AHO]
 - Useful for the first part of the course
- Other books
 - Cooper & Torczon. “Engineering a Compiler”
 - Hunter et al. “The essence of Compilers” (Prentice-Hall)
 - Grune et al. “Modern Compiler Design” (Wiley)
- Additional materials: notes, papers, technical reports



Students' Information

- Are you enrolled on ISIS (isis.tu-berlin.de)?
 - If not, do it!
- How many EIT students?
- How many Erasmus students?
- How many international exchange students?
- How many are enrolled to another German University (Humboldt U., Freie U., ...)?
- How many of you **cannot** be listed on QISPOS?

Important: Plagiarism NOT Allowed

- Exercises must contain original solutions only
- TU Berlin has strong rules against plagiarism
 - Fak. IV rules (German) <https://www.eecs.tu-berlin.de/fileadmin/f4/fkIVdokumente/plagiate.pdf>
 - What is plagiarism (English)
<https://www.ox.ac.uk/students/academic/guidance/skills/plagiarism?wssl=1>
 - Important: if you copy an exercise, you fail **the whole course**
- Also, you **cannot**
 - Upload your solution on the internet or share it on social networks
 - E.g., on GitHub



Syllabus

- Introduction
- Lexical analysis
 - Flex
- Syntax analysis
 - Bison
- Semantic analysis
- Intermediate representation
- Static Single Assignment (SSA)
- Dataflow analysis
 - LLVM analysis and optimization
- Code generation & runtime
- Register allocation
- Instruction scheduling
- Optimizations
- Compilation for embedded systems
 - LLVM code generation

Learning Outcomes

- A student successfully completing this course should be able to
 - understand the principles governing all phases of the compilation process
 - understand the role of each of the basic components of a standard compiler
 - show awareness of the problems of and methods and techniques applied to each phase of the compilation process
 - apply standard techniques to solve basic problems that arise in compiler construction
 - understating basic compiler optimizations and its implementation on real compiler (e.g., LLVM)
 - be a **better programmer**

What is a Compiler?

BEAR FACTS

by Burke



He always liked to open the Annual Programmer's Conference with a joke in binary.

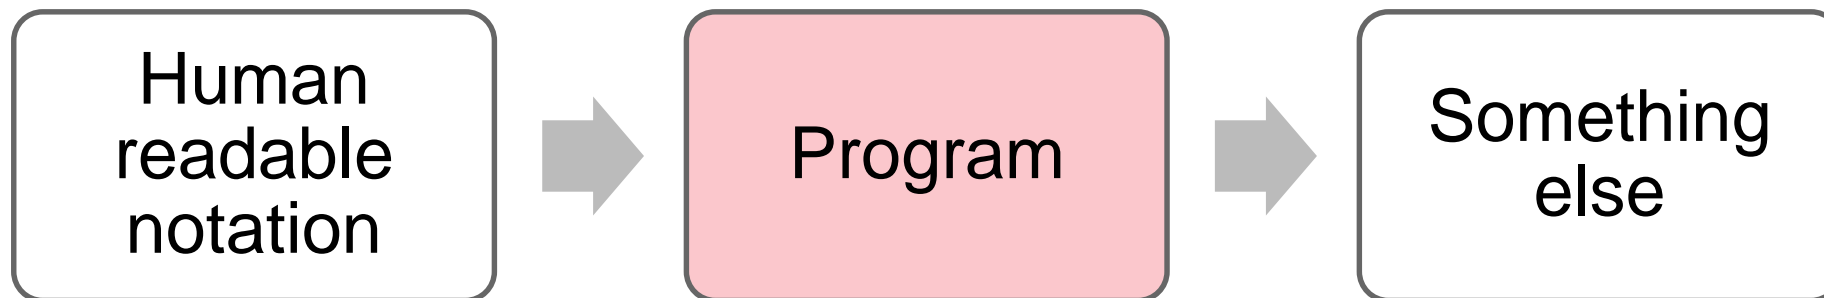
What is a Compiler?

“a computer program that translates a program written in a high-level language into another language, usually machine language”

➤ source: <http://dictionary.reference.com>

- German translations (from <http://dict.leo.org>)
 - Der Kompilierer
 - Der Übersetzer
 - Very interesting translation
 - Literally: **translator**

More Generally

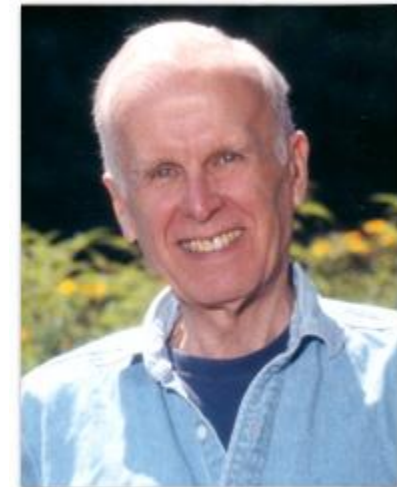


Historical Notes: Programming Languages

- Machine Languages
- 2nd generation: Assembly Languages – early 1950s
- 3rd generation: High-Level Languages – later 1950s
 - Fortran, ALGOL, COBOL
 - More recently C, C++, C#, Java, BASIC and Pascal,
- 4th generation higher level languages – 1970-1990
 - Also including domain specific languages (DSL)
 - SQL, Postscript, Python, Ruby, and Perl
- 5th generation languages: constraint-based, logic programming languages and some declarative languages
 - Prolog, OPS5, Mercury

Historical Notes: Fortran

- Fortran (Formula Translating System)
 - first, widely used **high-level programming language**
 - by John Warner Backus, IBM
 - in 1954 Backus assembled a team to define and develop Fortran for the IBM 704 computer
 - Backus also contributed to ALGOL58 and 60 and the BNF (Backus-Naur Form)
 - Turing Award, 1977



How many do you know?

javacc Cetus Polly/LLVM NESL (CMU)

gcc/g++

Rose Patus Charm (Illinois) Erlang

IBM xlc/xlcpp

PGI OpenMP compilers Sequoia (Stanford)

ARM armcc Pluto

ispc Insieme HPCS Chapel (Cray)

Intel icc/icpc HMPP

LLVM/Clang HPCS Fortress (Sun) Borealis (Brown)

NVIDIA cudacc

OpenCL compilers (Intel, AMD, NVIDIA, ...) Cilk (MIT/Intel) OpenACC

Definitions

- What is a **compiler**?

a program that accepts as input a program text in a certain language and produces as output a program text in another language, while preserving the meaning of that text [Grune et al., 2000]

a program that reads a program written in one language (source language) and translates it into an equivalent program in another language (target language) [AHO]

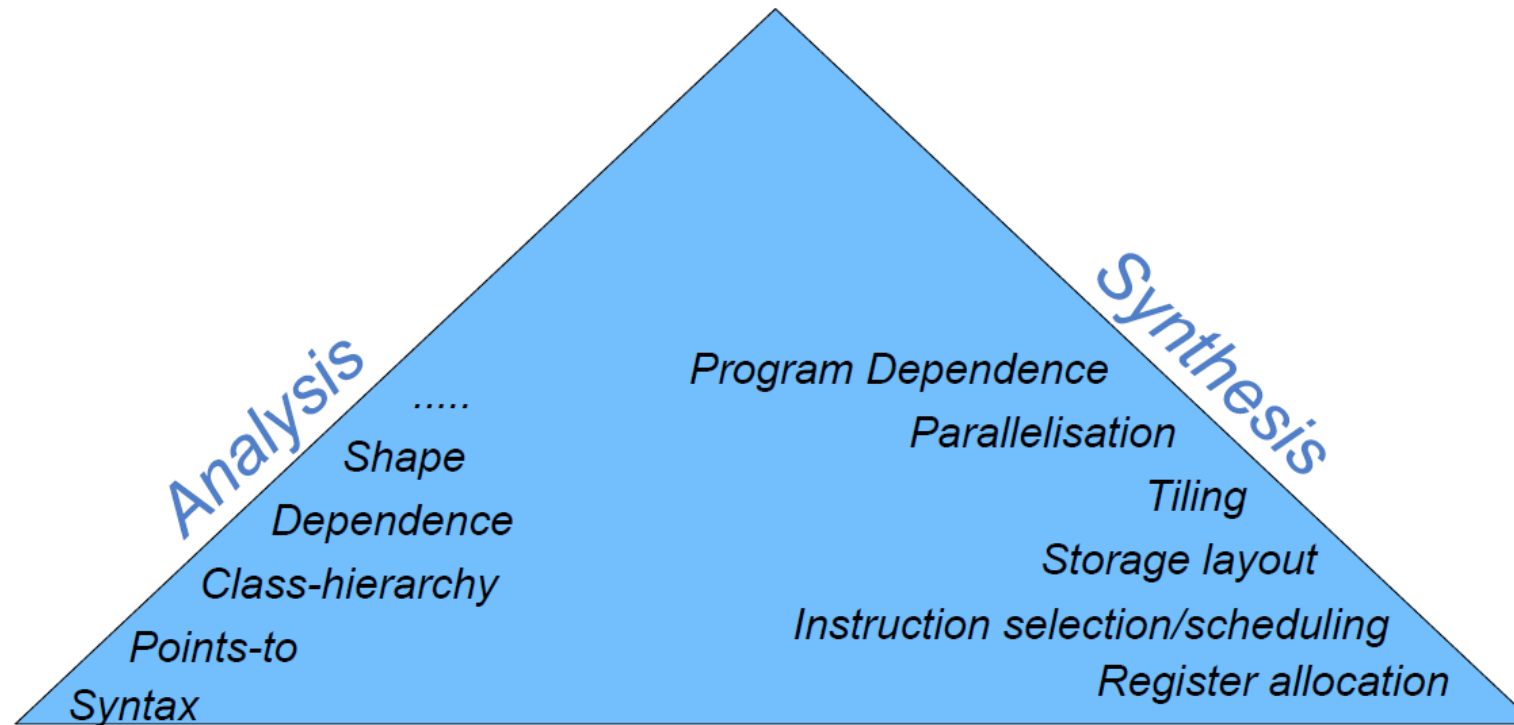
➤ key: ability to extract properties of a source program (**analysis**) and transform it to construct a target program (**synthesis**)

- What is an **interpreter**?

➤ a program that reads a source program and produces the results of **executing** this source

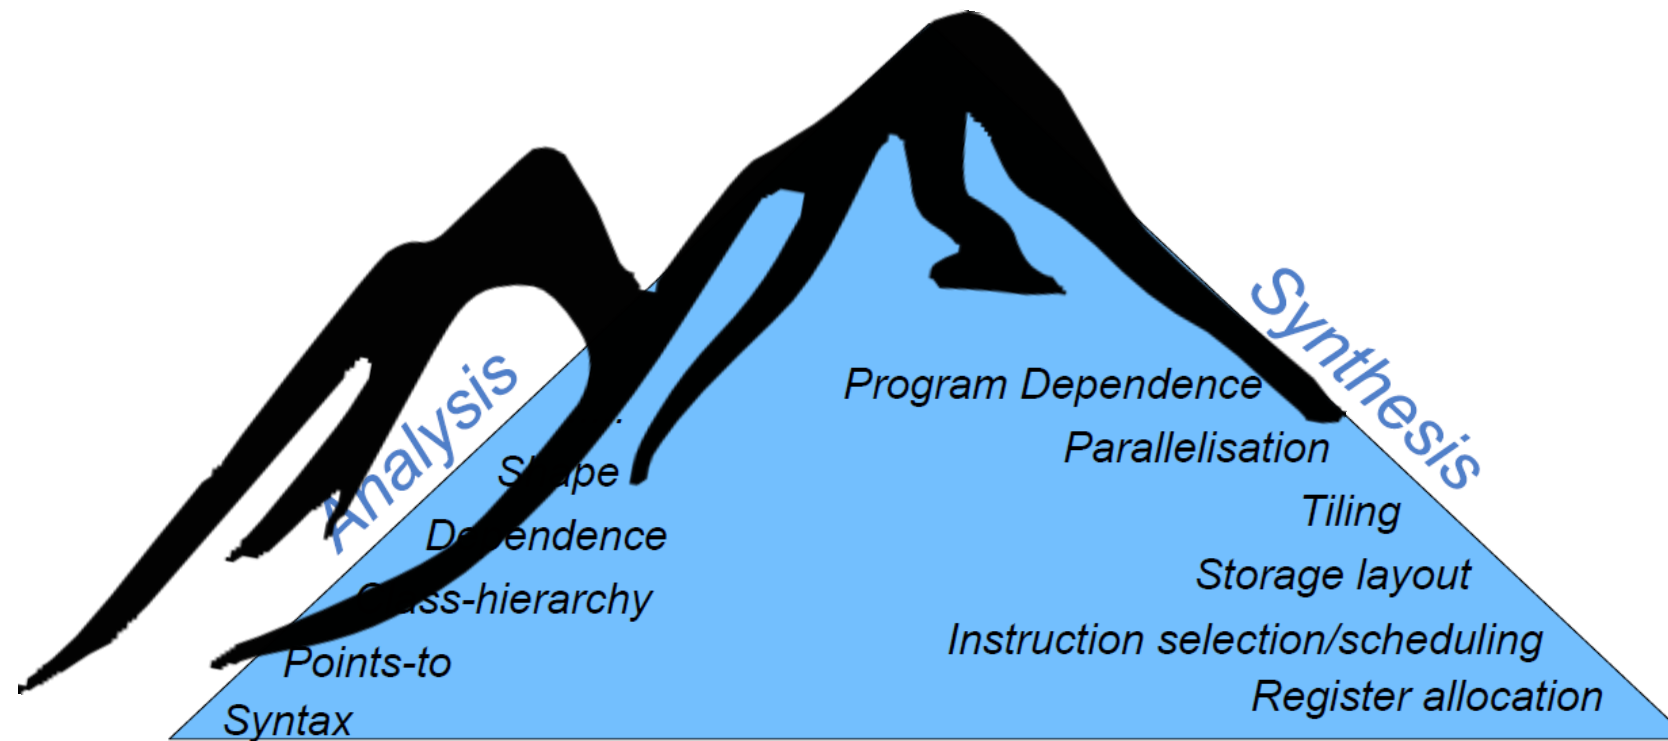
➤ an interpreter directly executes, i.e. performs, instructions written in a programming language, without previously compiling them into a machine language program

Analysis & Synthesis



Courtesy of Paul Kelly, Imperial College London

Analysis & Synthesis



Courtesy of Paul Kelly, Imperial College London

Questions

- Is it compiled or interpreted?
 - C
 - Lisp
 - Java
 - PHP
 - Latex
 - Ghostview
 - source-to-source C compiler
 - High Performance Fortran (HPF)

Example 1: Source-to-source Compilers

- Also called transpiler or transcompiler
- Typically C-to-C
- Code transformation at source level
 - examples: automatic parallelization, data layout transformations, ...
- High-level intermediate representation
 - we will see this in the next lectures (Intermediate Representations)
- Examples
 - Rose, Insieme, Pluto, Cetus, ...

Example 2: Java JIT Compiler

- Java compiler
 - the output is a class file (.class)
 - `javac` (Oracle), `gcj` (GNU Compiler for Java), `ECJ` (Eclipse for Java)
 - platform-neutral Java bytecode
 - there are also compilers that emit optimized native machine code for a particular hardware/operating system combination
- Most Java-to-bytecode compilers do little optimization, leaving this to the JRE (Java Runtime) at runtime
- **Just-in-time** (JIT) compilation
 - the Java virtual machine (JVM) loads the class files and either **interprets** the bytecode or **just-in-time** compiles it to machine code and then possibly optimizes it using dynamic compilation.
 - Interaction between JVM and Java compilers specified in JSR 199

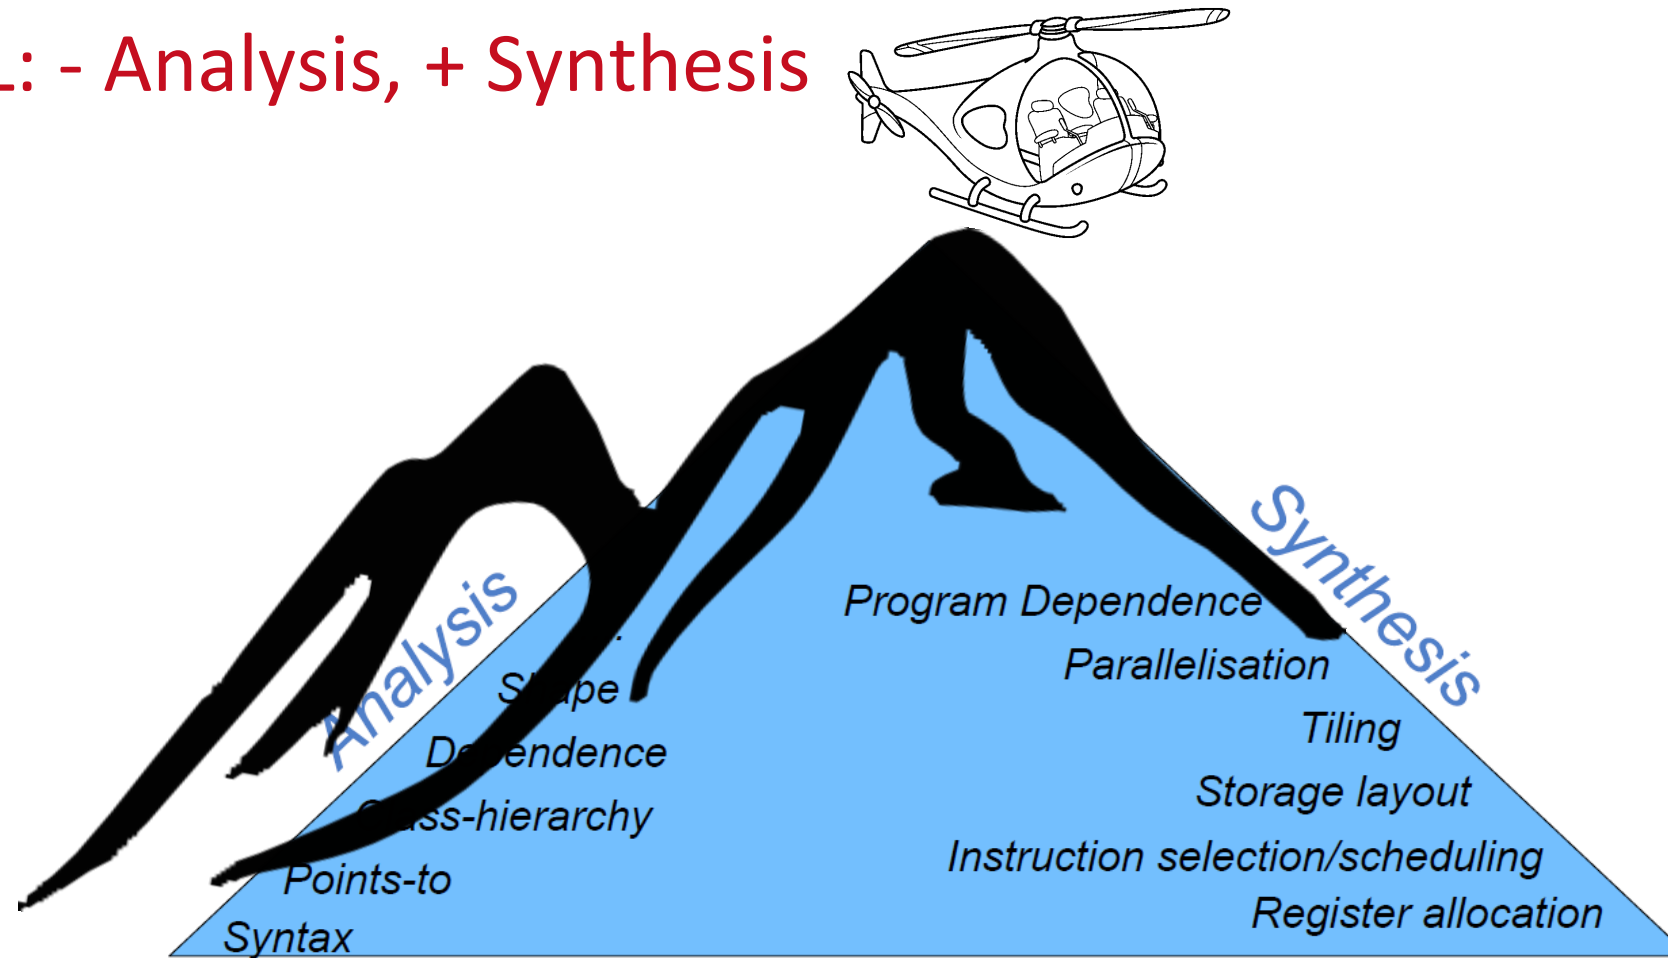
Example 3: Domain-Specific Languages

- Input is a domain-specific language (**DSL**)
 - a language with domain-specific construct and constraints
- Assumptions (and restriction) on the input
 - analysis simpler
 - optimization is relatively easier
 - Domain-specific optimization
- Examples
 - OpenGL Shading Language
 - Halide for image processing
 - SQL for database

Example: simple GLSL
fragment shader

```
varying vec3 N;  
varying vec3 v;  
  
void main(void){  
    vec3 L = normalize(gl_LightSource[0].position.xyz-v);  
    vec4 Idiff = gl_FrontLightProduct[0].diffuse *  
max(dot(N,L), 0.0);  
    Idiff = clamp(Idiff, 0.0, 1.0);  
    gl_FragColor = Idiff;  
}
```

DSL: - Analysis, + Synthesis



Courtesy of Paul Kelly, Imperial College London

Example 4: Parallelizing Compilers

- Input is sequential code
- Output is parallel code, i.e., expose some kind of parallelism
- Typically, parallelism is extracted from loops
 - advanced analysis, e.g., using the polyhedral model
 - `#pragma` notations to help the compiler job
- Sometime parallelizing compilers enhance parallelization
 - E.g., from shared memory parallel code to distributed or heterogeneous systems
- Form of parallelism
 - automatic vectorization (SIMD instructions), e.g., by `gcc`, `llvm` and `icc`
 - multi-threading (pthread), e.g., by Rose, Pluto, Insieme, LLVM-Polly
 - distributed memory (typically MPI)

An Advanced Example: Automatic Parallelization

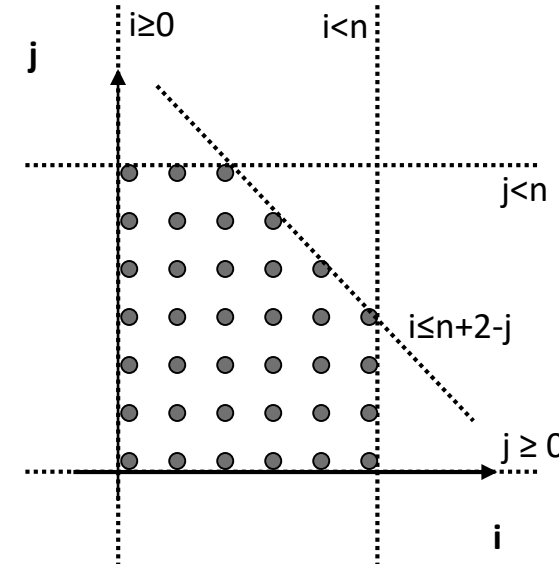
- How can you make this code parallel?

```
for(int i=0; i<n; i++)  
    for(int j=0; j<n; j++)  
        if(i <= n+2-j)  
            b[j] = b[j] + a[i];
```

- Are iterations independent?
- Suppose you have four processors
 - How would you represent (and distribute) loop iterations between processors?

Automatic Parallelization & Polyhedral Model

```
for(int i=0; i<n; i++)
  for(int j=0; j<n; j++)
    if(i <= n+2-j)
      (s)      b[j] = b[j] + a[i];
```



Polyhedron for n=6



$$\begin{bmatrix} 1 & 0 \\ -1 & 0 \\ 0 & 1 \\ 0 & -1 \\ -1 & -1 \end{bmatrix} \begin{pmatrix} i \\ j \end{pmatrix} + \begin{pmatrix} 0 \\ n-1 \\ 0 \\ n-1 \\ -n-2 \end{pmatrix} \geq \vec{0}$$

Iteration domain of S

i j n constant

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 0 & 1 & -1 \\ 0 & 1 & 0 & 0 \\ 0 & -1 & 1 & -1 \\ -1 & -1 & -1 & -2 \end{bmatrix} \begin{pmatrix} i \\ j \\ n \\ 1 \end{pmatrix} \geq \vec{0}$$

Iteration domain with homogenous coord.



Qualities of a Good Compiler

- What qualities would you like in a compiler?
 - generates correct code (first and foremost!)
 - generates fast code
 - conforms to the specifications of the input language
 - copes with essentially arbitrary input size, variables, etc.
 - compilation time (linearly) proportional to size of source
 - good diagnostics
 - consistent optimizations
 - works well with the debugger

Principles of Compilation

- The compiler must
 - preserve the meaning of the program being compiled
 - “improve” the source code in some way
- Other issues (depending on the setting)
 - speed (of compiled code)
 - space (size of compiled code), energy consumption
 - feedback, latency (information provided to the user)
 - debugging (transformations obscure the relationship source code vs target)
 - compilation time efficiency (fast or slow compiler?)

Uses of Compilers

- Simply translation of high-level program to object code
 - program translation: binary translation, hardware synthesis, ...
- Optimizations
 - improve program performance, take into account hardware
 - automatic parallelization
- Performance instrumentation
 - example: `-pg` option of `gcc`
- Interpreters
 - Perl, bash, ...
- Software productivity tools
 - debugging aids, e.g. `purify`
- Security
 - Java VM uses compiler analysis to prove “safety” of Java code.
- Web-browsers (Javascript and HTML), text formatters, just-in-time compilation for Java, power management, global distributed computing, ...

*Ability to extract properties of a source program (**analysis**) and transform it to construct a target program (**synthesis**)*



Questions

- Difference between compiler and interpreter
- What is a source-to-source compiler?
- What is a parallelizing compiler?
- What is a Domain Specific Language?

Programming Languages

L_1

L_2

\vdots

L_m

Target Architecture

T_1

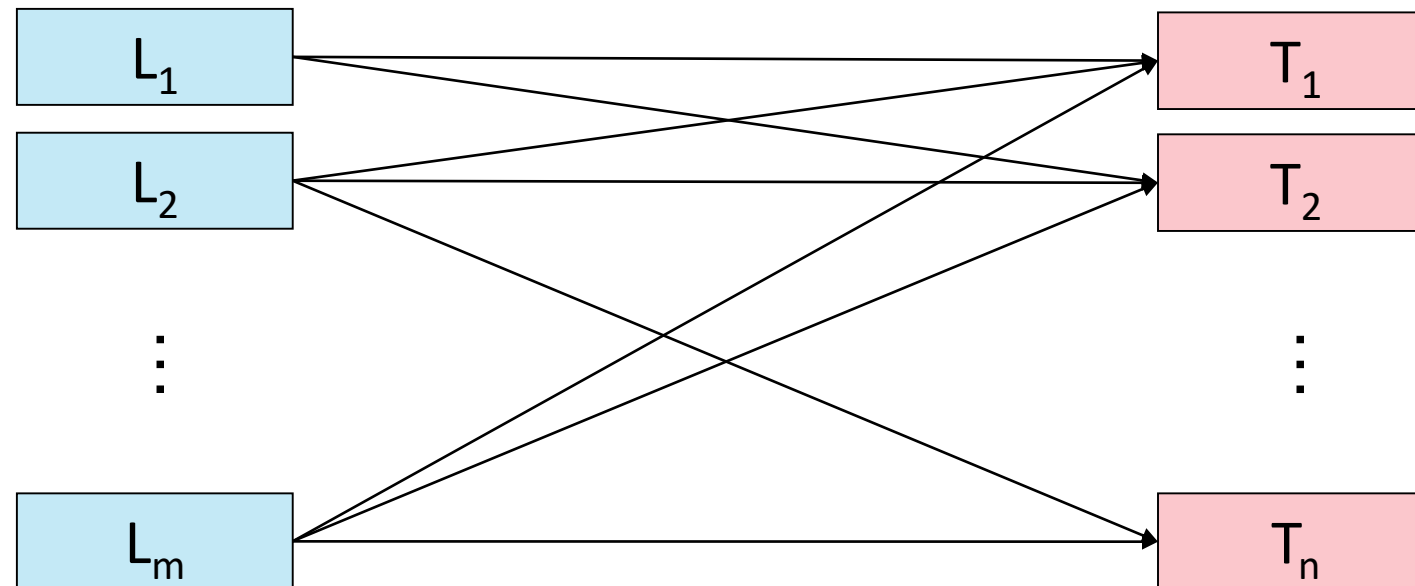
T_2

\vdots

T_n

Programming Languages

Target Architecture



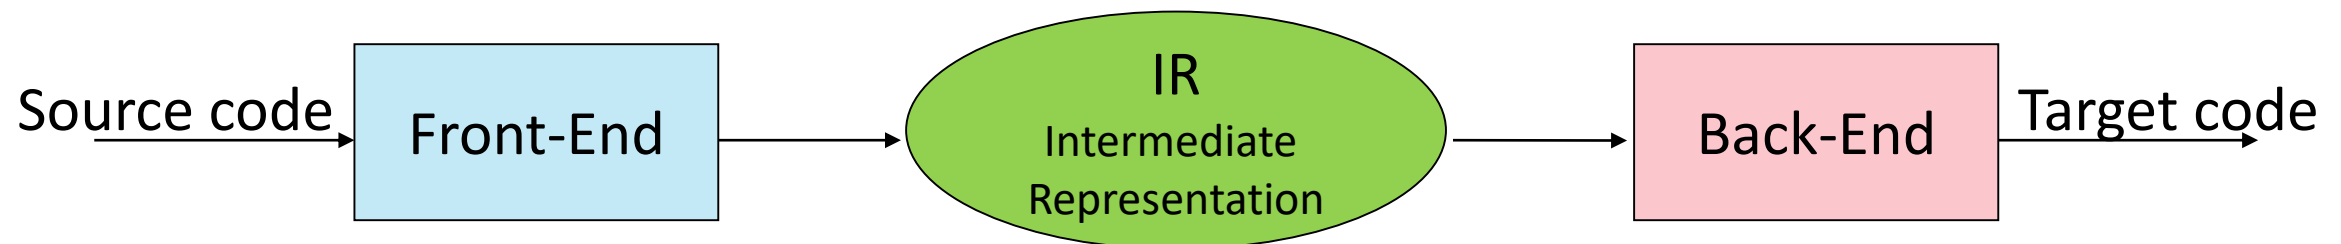
General Structure of a Compiler

Front-end performs the analysis of the source language

- recognizes legal and illegal programs and reports errors
- understands the input program and collects its semantics in an **Intermediate Representation** (IR)
- produces IR and shapes the code for the back-end

Back-end does the target language synthesis

- chooses instructions to implement each IR operation
- translates IR into target code
- needs to conform with system interfaces
- automation has been less successful

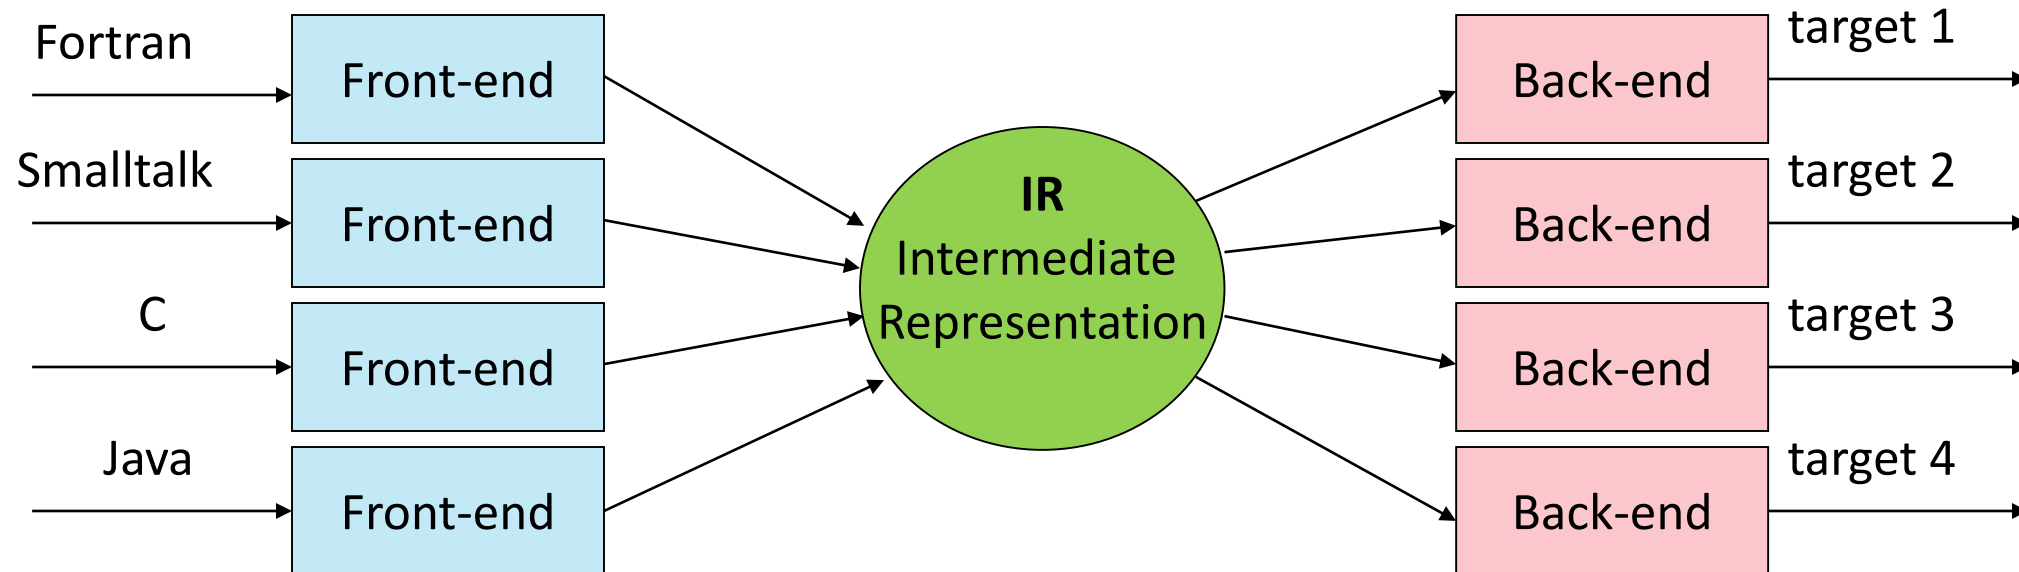




Questions

- What is the implication of this separation (front-end: analysis; back-end: synthesis) in building a compiler for a new language?
- And for a new target architecture?

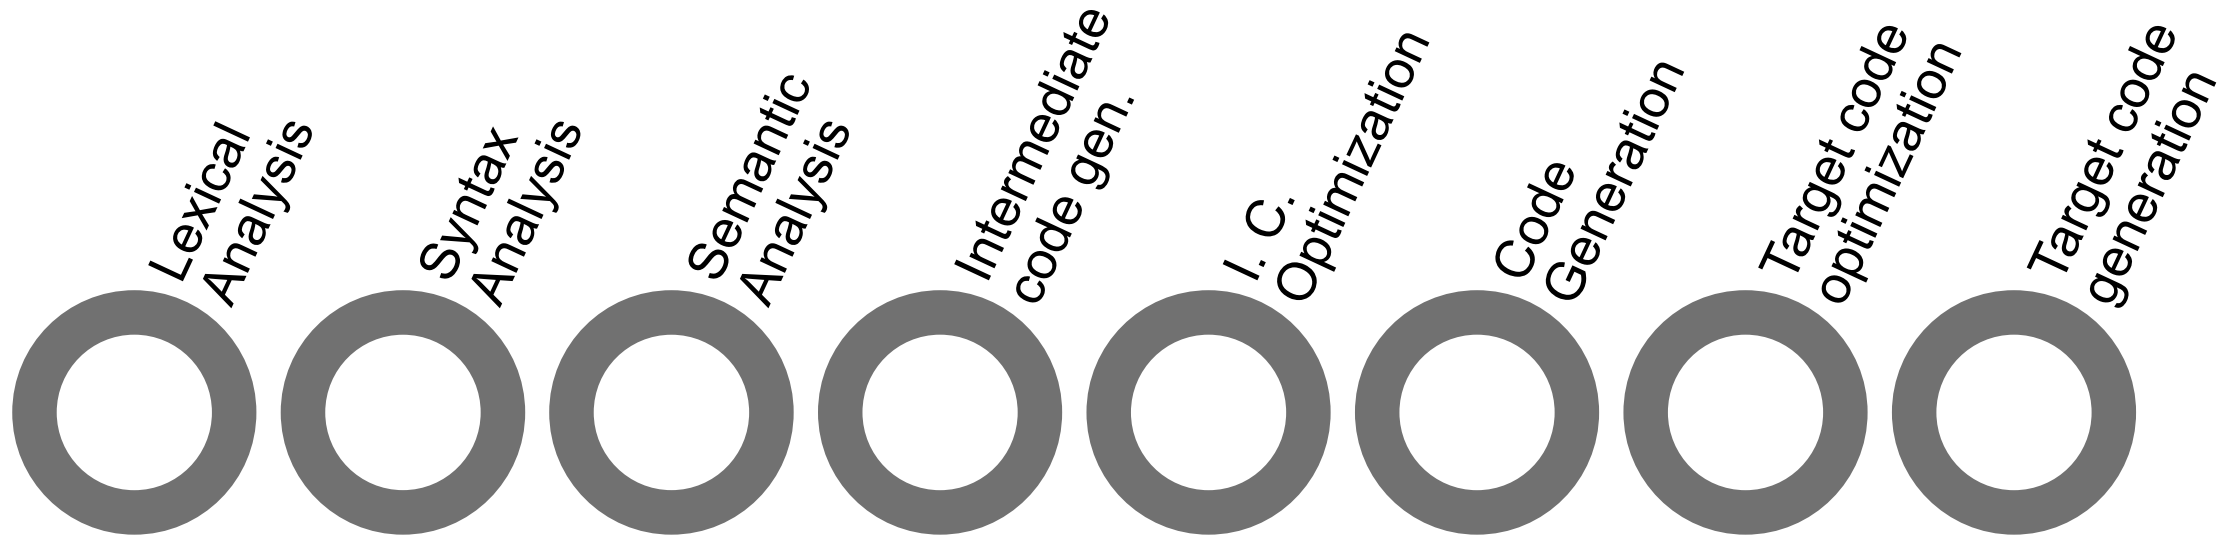
Answer: $m \times n$ compilers with $m+n$ components!



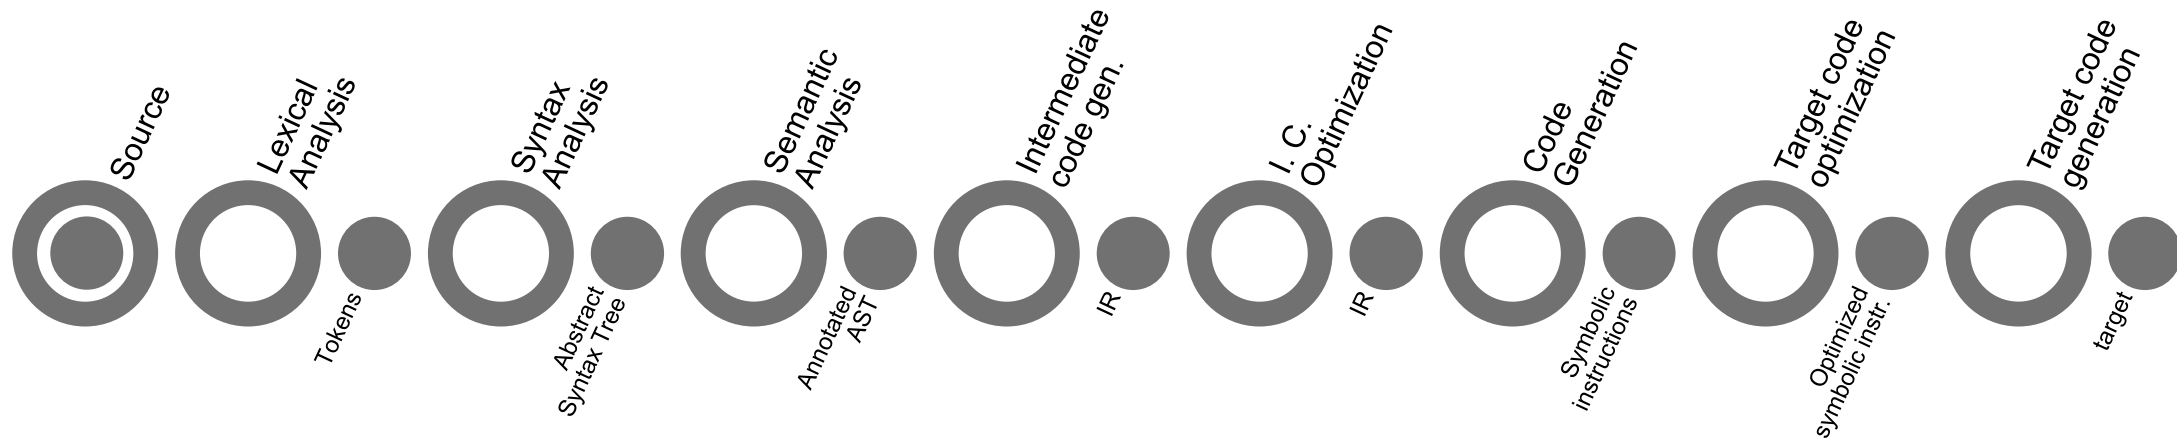
- All language-specific knowledge must be encoded in the front-end
- All target-specific knowledge must be encoded in the back-end

This strict separation
is not free of charge!

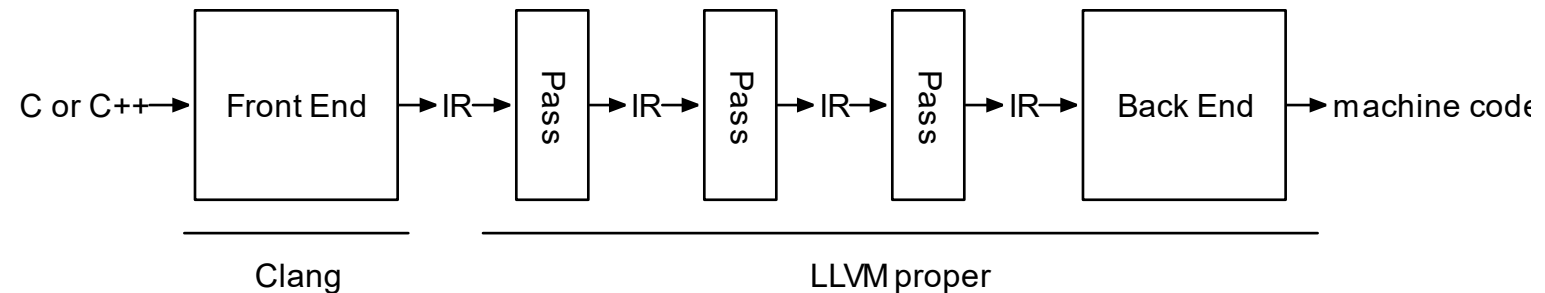
General Compiler Structure



General Compiler Structure



Example:
the **LLVM** compiler



Lexical Analysis

- Reads characters in the source program and groups them into words (basic unit of syntax)
- Produces words and recognizes what sort they are
- The output is called token and is a pair of the form `<type, lexeme>` or `<token_class, attribute>`
 - E.g.: **a=b+c** becomes `<id, a>` `<=, >` `<id, b>` `<+, >` `<id, c>`
- Needs to record each id attribute: keep a symbol table
 - Lexical analysis eliminates white space, etc
- Speed is important - use a specialized tool
 - **Flex**: a tool for generating scanners: programs which recognize lexical patterns in text

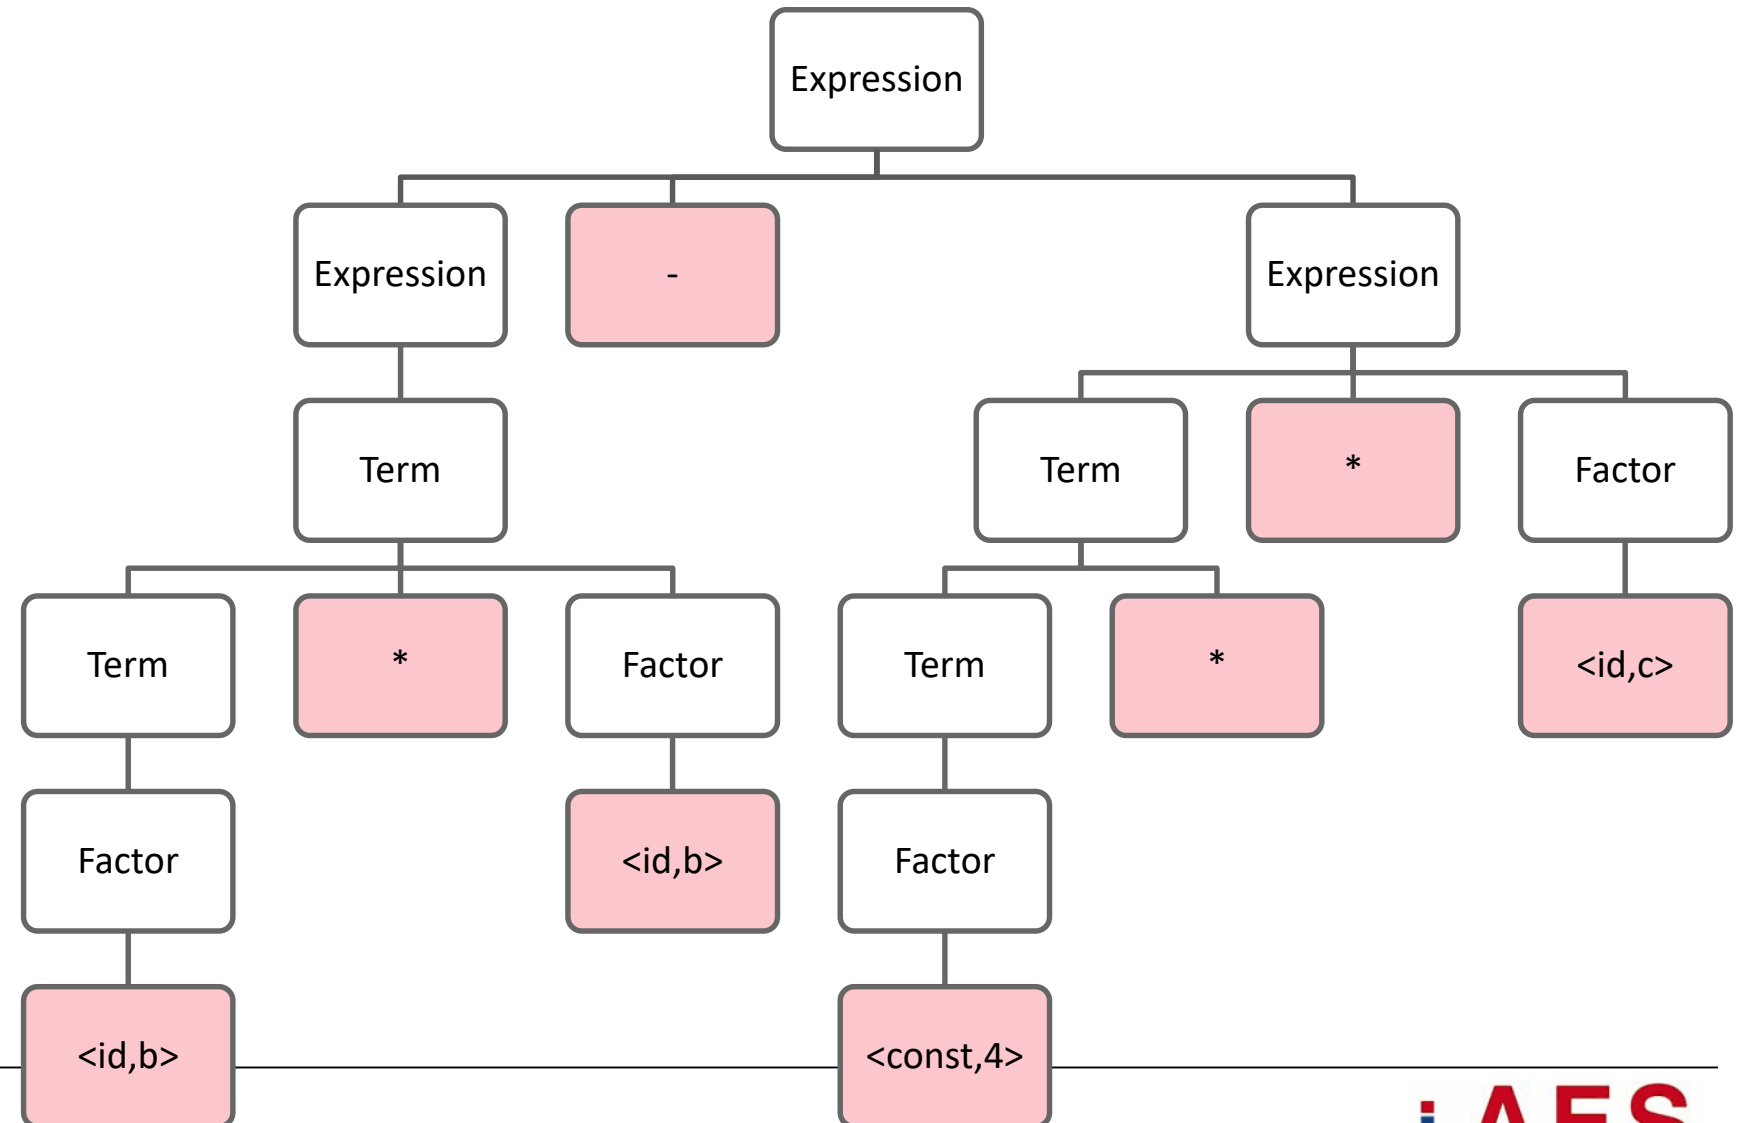
Syntax Analysis (Parsing)

- Imposes a hierarchical structure on the token stream
- This hierarchical structure is usually expressed by recursive rules
- Context-free grammars formalise these recursive rules and guide syntax analysis
- Example
 - A grammar defining simple algebraic expressions

```
expression -> expression '+' term | expression '-' term | term
term -> term '*' factor | term '/' factor | factor
factor -> identifier | constant | '(' expression ')'
```

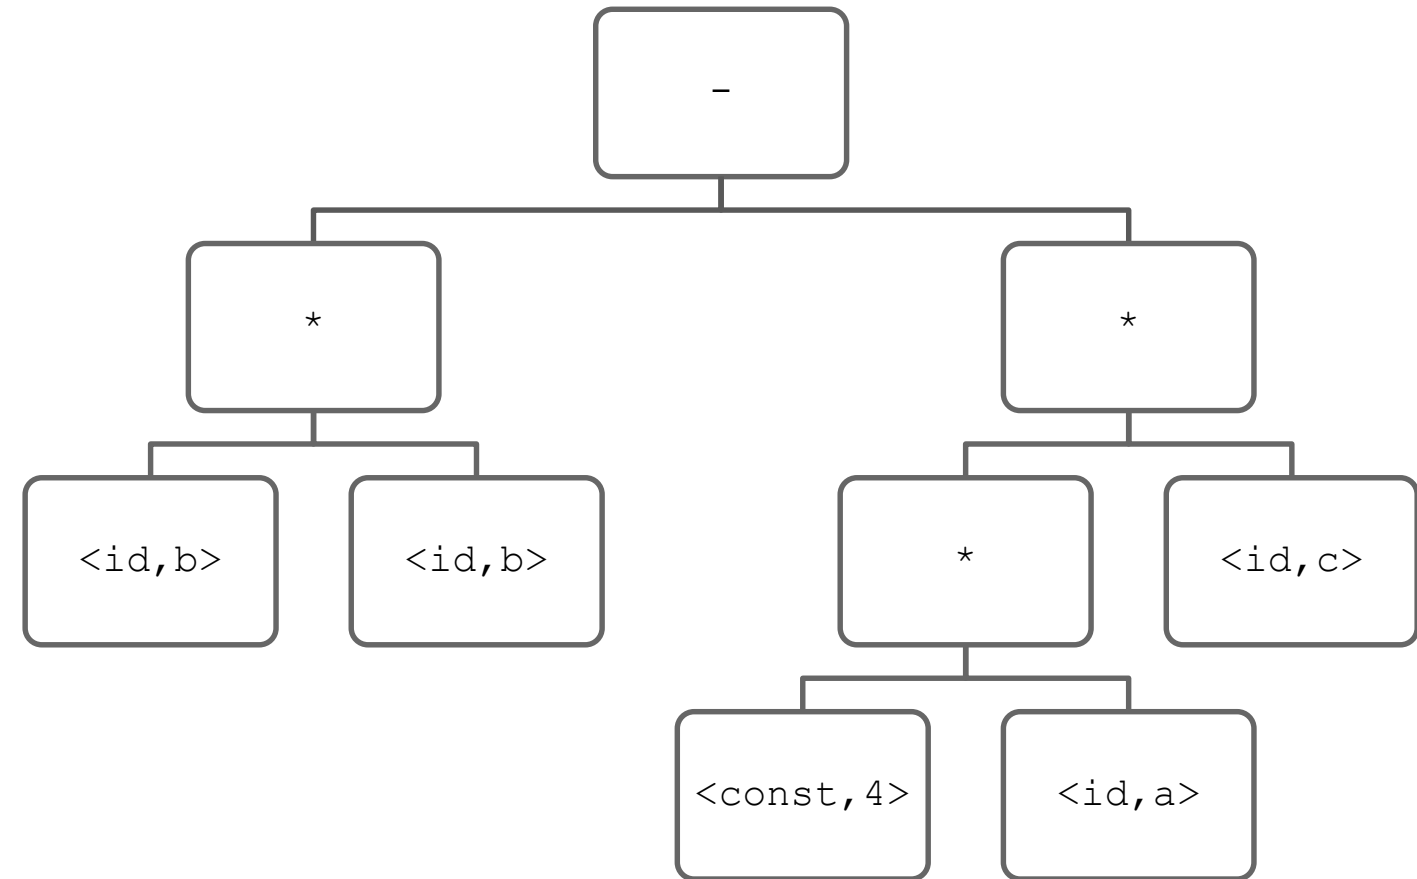

Parsing

- Parse tree for $b*b-4*a*c$



Abstract Syntax Tree (AST)

- AST for $b*b-4*a*c$
- An **Abstract Syntax Tree (AST)** is a more useful data structure for internal representation
- It is a compressed version of the parse tree (summary of grammatical structure without details about its derivation)
- ASTs are a form of IR



Semantic Analysis (Context Handling)

- Collects context (semantic) information, checks for semantic errors, and annotates nodes of the tree with the results
- Examples
 - type checking: report error if an operator is applied to an incompatible operand
 - check flow-of-controls
 - uniqueness or name-related checks

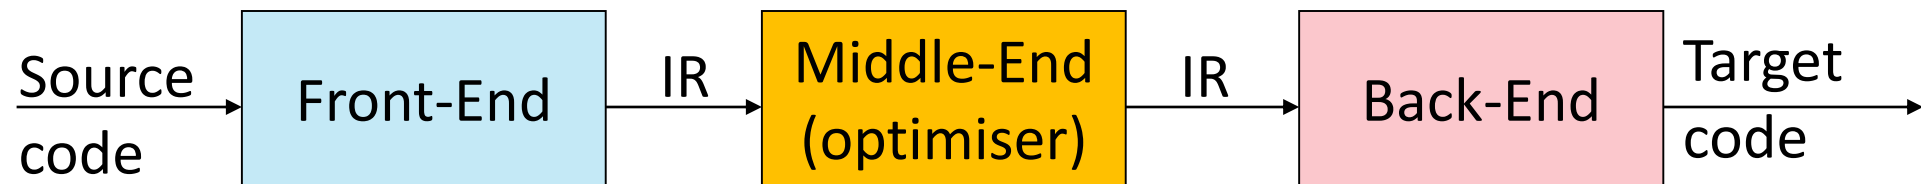
Intermediate Code Generation

- Translate language-specific constructs in the AST into more general constructs
- A criterion for the level of “generality”: it should be straightforward to generate the target code from the intermediate representation chosen
- Example of a form of IR (3-address code)

```
tmp1 = 4  
tmp2 = tmp1*a  
tmp3 = tmp2*c  
tmp4 = b*b  
tmp5 = tmp4-tmp3
```

Code Optimization

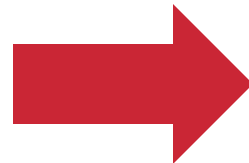
- The goal is to improve the intermediate code and, thus, the effectiveness of code generation and the performance of the target code
- Optimizations can range from trivial (e.g., constant folding) to highly sophisticated (e.g., inlining)
 - Example: replace the first two statements in the example of the previous slide with: $\text{tmp2} = 4 * a$
- Modern compilers perform such a range of optimizations, that one could argue for:



Optimizations

- Example: Dead Code Elimination (**DCE**)

```
int global;  
void f ()  
{  
    int i;  
    i = 1;  
    global = 1;  
    global = 2;  
    return;  
    global = 3;  
}
```



```
int global;  
void f ()  
{  
  
    global = 2;  
    return;  
}
```

Code Generation

- Map the AST onto a linear list of target machine instructions in a symbolic form
 - instruction selection: a pattern matching problem
 - register allocation: each value should be in a register when it is used (but there is only a limited number): NP-Complete problem
 - instruction scheduling: take advantage of multiple functional units: NP-Complete problem
- Target, machine-specific properties may be used to optimize the code
- Finally, machine code and associated information required by the Operating System are generated



Questions

- List the phases of a compiler
- What is the front-end?
- What is the back-end?
- What is the IR?
- Why, in a compiler infrastructure, the front-end is typically separated from the back-end?

Summary

- Compiler vs interpreter
- Frontend, backend, IR
- Structure of a compiler and compilation phases
- Source-to-source, DSL
- Recommended reading:
 - [ALSU] 1.1, , 1.2, 1.3, 1.5
 - JSRs: Java Specification Requests - JSR 199: Java Compiler API
 - <https://jcp.org/en/jsr/detail?id=199>
 - J. W. Backus, *Can Programming be Liberated from the von Neumann Style?*
 - <https://www.cs.cmu.edu/~crary/819-f09/Backus78.pdf>
 - H. Massalin, *Superoptimizer: A look at the smallest program*
 - In ASPLOS II, pages 122-126, Los Alamitos, CA, USA, 1987. IEEE Computer Society Press.

